

Jagan Guru Teja Ponugupati

Data Engineer

jagangurutejajob@gmail.com | +91 809-608-3180

PROFESSIONAL SUMMARY

Data Engineer with 3 years and 6 months of experience designing and developing scalable ETL/ELT pipelines for enterprise data integration, cloud migration, and analytics. Hands-on expertise in Azure Data Factory, Microsoft Fabric (Lakehouse, Warehouse, Notebooks), SQL, and Python/PySpark, with experience implementing Medallion Architecture, metadata-driven pipelines, incremental loading, Delta Lake, SCD Type 1 & 2, and dimensional modeling. Experienced in migrating legacy SQL Server and SSAS workloads to Azure and Microsoft Fabric, building data warehouse solutions and Power BI-ready datasets, and managing CI/CD using Azure DevOps and Git.

TECHNICAL SKILLS

Cloud & Data Engineering: Azure Data Factory (ADF), Microsoft Fabric, Azure Data Lake Storage Gen2, Azure Blob Storage, Azure Function Apps, Azure Key Vault, SSIS

Microsoft Fabric: Data Pipelines, Lakehouse, Warehouse, PySpark Notebooks, Semantic Models, OneLake (Shortcuts), Medallion Architecture

ETL & Data Processing: ETL/ELT, Metadata-Driven Pipelines, Incremental Loading, Delta Lake, Delta MERGE, SCD Type 1 & 2, Data Validation & Reconciliation

Languages: SQL, Python, PySpark, SOQL, MQL

Databases & Data Warehousing: SQL Server, Azure SQL Database, MongoDB, Oracle, SSMS, SSAS, Data Warehousing, Dimensional Modeling

APIs & Integration: REST APIs, Salesforce CPQ, Azure Functions

DevOps & Source Control: Azure DevOps, CI/CD, Git, GitHub

BI & Analytics: Power BI

PROFESSIONAL EXPERIENCE

Data Engineer — Amnet Digital, Hyderabad

April 2023 – Present

CPQ Data Migration & Integration — Hardware, Software & Subscription Orders

Dec 2023 – Present | Client: CDW | Retail / Sales Operations (CPQ) | Data Engineer

Tech: MongoDB, Azure Data Factory, SQL Server, Salesforce CPQ, Azure Key Vault, Azure Blob Storage, Rest APIs, GitHub, Azure DevOps

- Built end-to-end ETL pipelines in Azure Data Factory to extract, transform, and load hardware, software, and subscription order data from MongoDB into Salesforce CPQ, using Azure SQL Server as the staging and transformation layer.
- Wrote SQL Server stored procedures to handle data transformation, validation, and error handling, logging error rows into a dedicated error table to catch and diagnose issues from the source data.
- Implemented copy activity logging to track success and failure at the activity level, and pipeline-level logging using a master-child pipeline structure for centralized tracking; added email alerts for pipeline success and failure to enable quick issue detection.
- Set up real-time, event-driven data sync using Azure Blob Storage triggers alongside scheduled and manual reprocessing workflows, resolving deadlocks between environments and keeping Commercial, Canada, and Government environments reliably in sync.
- Built validation and reconciliation checks to catch discrepancies and ensure data stayed accurate and consistent across the source, staging, and Salesforce CPQ systems.
- Used Azure DevOps CI/CD pipelines for deployment and GitHub for version control, ensuring consistent and traceable releases across environments.
- Worked closely with data, integration, and business teams to resolve migration issues, confirm requirements, and adapt to changing schemas and CPQ integration needs.

Enterprise Data Services & Analytics: Legacy-to-Cloud & Fabric Migration

2023 – Present | Client: CDW | Data Engineer

Tech: SSMS, SSIS, SSAS, Azure Data Factory, Azure Blob Storage, Azure Key Vault, Azure Data Lake Storage Gen2, REST APIs, Azure Function Apps, Microsoft Fabric (Warehouse, PySpark Notebooks), Power BI, Azure DevOps, Git

- Developed Azure Data Factory (ADF) ETL pipelines to migrate tables and stored procedures from on-premises SQL Server to Azure, centralizing data from multiple source systems while preserving existing business logic.
- Built parameterized, metadata-driven ADF pipelines using Copy Activity, Lookup, Stored Procedures, and ForEach to ingest records from SQL Server, REST APIs.
- Implemented incremental loading, Delta MERGE, and SCD Type 1 & 2 in ADF to maintain accurate and up-to-date analytics data.
- Migrated SSAS Cube logic to Microsoft Fabric, re-engineering transformations using PySpark and SQL, designing Star Schema-modeled Fabric Warehouse tables for Power BI, and decommissioning the legacy SSAS environment.
- Implemented data validation, reconciliation, and monitoring/audit logging to ensure reliable, high-quality pipelines across the migration.
- Validated migrated data using a reusable Python automation script, ensuring accuracy and consistency between source and target systems post-migration.
- Managed Git-based source control and CI/CD deployments using Azure DevOps across development, testing, and production environments.

CERTIFICATIONS

Microsoft Certified: Fabric Analytics Engineer Associate (DP - 600)

Earned Nov 2024 | Valid through Nov 2027

Microsoft Certified: Fabric Data Engineer Associate (DP - 700)

Earned Aug 2026 | Valid through Aug 2027

EDUCATION

Bachelor of Technology — Sir CR Reddy College of Engineering, Eluru

2019 – 2022 | CGPA: 7.35